



Writing Faster Python

Roger Hatfull


15th Annual Symposium for Graduate Physics Research
GPSA, University of Alberta

<https://tinyurl.com/wfpgpsa2024>

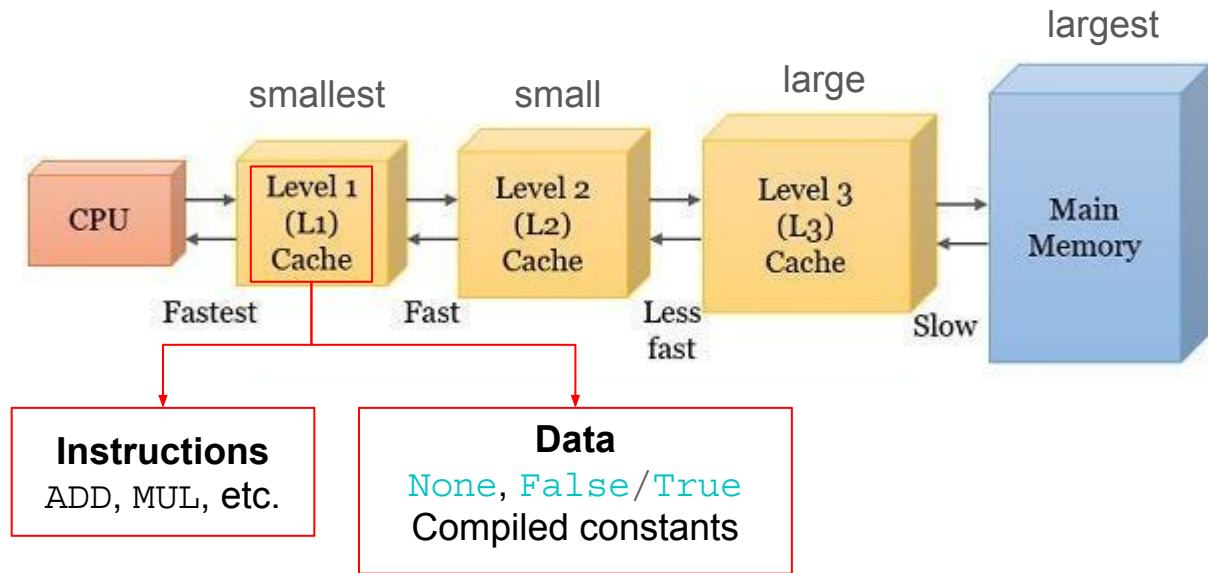
<https://github.com/hatfullr/writing-faster-python>



~20 TB data
~300 CPU years
~50 GPU years

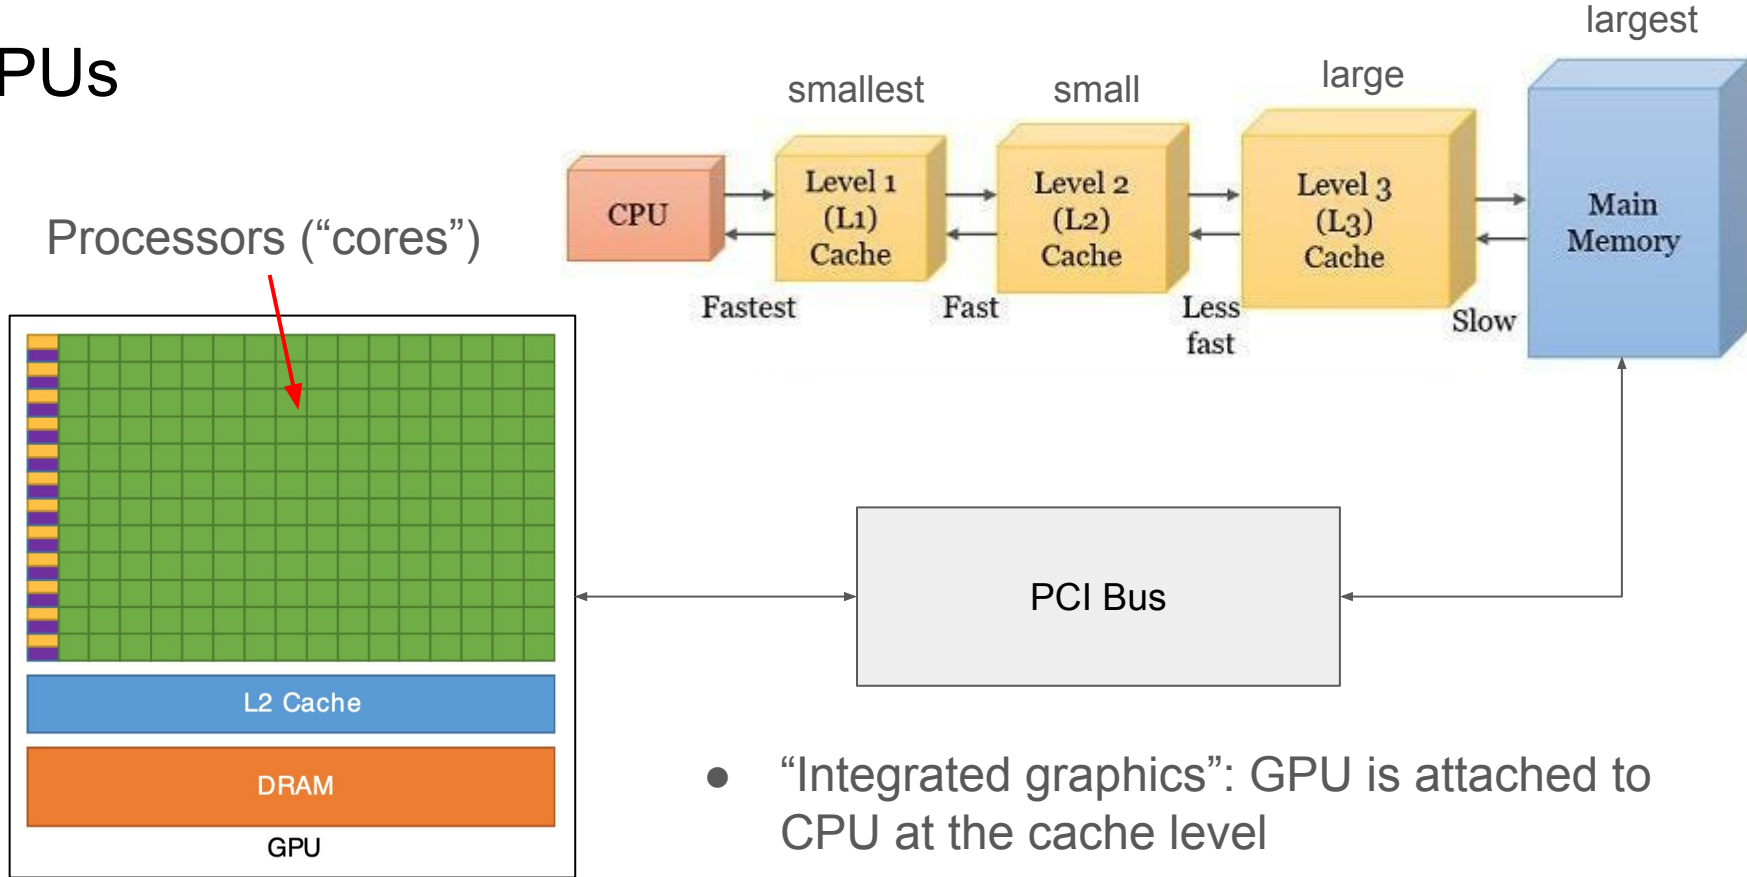


CPUs



- (Some) manufacturer tricks:
 - “Cache lines”: data from caches are loaded 64 bytes at a time
 - “Pre-fetching”: Predict and load to low-level caches before needed
 - Scripts run faster on consecutive executions: data is being cached

GPUs



Lowest-hanging fruit

- **Programmer time > execution time**
- Make shortest syntax a habit
- Modern computations are typically memory-bound
 - Performing calculations can be much faster than storing calculations

scripts/syntax/list.py

```
0)    0.528331 seconds    "for i in range(len(a)): b += [a[i]] "  
1)    0.485454 seconds    "for i, ai in enumerate(a): b += [ai] "  
2)    0.479977 seconds    "for i in range(len(a)): b[i] = a[i] "  
3)    0.464579 seconds    "for i, ai in enumerate(a): b[i] = ai "  
4)    0.424565 seconds    "for i in range(len(a)): b.append(a[i]) "  
5)    0.398427 seconds    "for i, ai in enumerate(a): b.append(ai) "  
6)    0.362043 seconds    "for ai in a: b += [ai] "  
7)    0.290544 seconds    "for ai in a: b.append(ai) "  
8)    0.279091 seconds    "b = [ai for ai in a] "  
9)    0.105188 seconds    "b = a.copy() "
```

(Some) coding practices

- Avoid indentations
 - Unless it makes the code easier to read
- Loops with conditionals:
 - Use `continue` and `break`

scripts/examples/practices.py

```
for i in range(3):  
    if i == 0:  
        print("Hello")  
        continue  
    if i == 1:  
        print("What's up?")  
        continue  
    print("Not much")
```

scripts/examples/practices.py

```
for i in range(3):  
    if i == 0:  
        print("Hello")  
    else:  
        if i == 1:  
            print("What's up?")  
        else:  
            print("Not much")
```

\$ python3 practices.py
Hello
What's up?
Not much

if condition: value = True
else: value = False ❌

value = condition ✅

Being careful

- What are pointers?
 - When the computer accesses a pointer in memory, it receives an instruction to access a different place in memory
- Python creates pointers sometimes
- If you aren't sure: `copy.deepcopy()`

scripts/examples/pointers.py

```
a = [False, True, False]
b = a
print(a)
b[1] = False
print(a)
```

```
$ python3 pointers.py
[False, True, False]
[False, False, False]
```

scripts/examples/pointers.py

```
import copy
a = [False, True, False]
b = copy.deepcopy(a)
print(a)
b[1] = False
print(a)
```

```
[False, True, False]
[False, True, False]
```

Memory efficiency

- Generators (“enumerators”)

- `yield`
- Evaluates expressions on-the-fly
- Saves memory space
- Allows to easily write complex algorithms

```
def find_files(directory):  
    results = []  
    for filename in os.listdir(directory):  
        path = os.path.join(directory, filename)  
        if os.path.isdir(path):  
            results += find_files_regular(path)  
            continue  
        else: results += [path]  
    return results
```

(snippet) `scripts/examples/generators.py`

```
def find_files(directory):  
    for filename in os.listdir(directory):  
        path = os.path.join(directory, filename)  
        if os.path.isdir(path): yield from find_files(path)  
        else: yield path
```


Faster NumPy

- Avoid “np. _____” at all costs

scripts/syntax/np.py

0)	4.605282	seconds	"for i, ai in enumerate(a): b = np.append(b, ai) "
1)	0.374851	seconds	"for i, ai in enumerate(a): b[i] = ai "
2)	0.103949	seconds	"np.copy(a) "
3)	0.058146	seconds	"a.copy() "
0)	2.729569	seconds	"np.sum(a) "
1)	2.181445	seconds	"for i, ai in enumerate(a): b += ai "
2)	1.691387	seconds	"sum(a) "
3)	1.342459	seconds	"a.sum() "
0)	1.332396	seconds	"np.power(a, 2) "
1)	0.541881	seconds	"pow(a, 2) "
2)	0.525106	seconds	"a**2 "
3)	0.509239	seconds	"a*a "

line_profiler



scripts/examples/line_profiler_example.py

https://github.com/pyutils/line_profiler

```
@profile
def func(): # Some long operation
    for i in range(int(1e6)):
        i**2

func()
```

```
$ kernprof -l line_profiler_example.py
Wrote profile results to line_profiler_example.py.lprof
$ python3 -m line_profiler line_profiler_example.py.lprof
Timer unit: 1e-06 s
```

```
Total time: 0.912989 s
File: line_profiler_example.py
Function: func at line 1
```

Line #	Hits	Time	Per Hit	% Time	Line Contents
1					@profile
2					def func(): # Some long
operation					
3	1000001	330973.0	0.3	36.3	for i in range(int(1e6)):
4	1000000	582016.0	0.6	63.7	i**2

Writing for the “user”

● YOU ARE THE USER

- Be kind to yourself!!
 - Function annotations
 - Doc strings
 - Code comments: “why” not “what”
- If not, you WILL LOSE TIME LATER
- Before writing, ask “is this something I might use a lot?”
- Most difficult task in programming is understanding someone else’s code
- You will be someone else in ~6 months
 - Your old code == someone else’s code

```
def useless(
    param1 : float,
    param2 : int,
    param3 : type(None) | str = None,
):
    """
    Describe the function. Doc strings can be used later for automatic
    documentation (see sphinx), and for remembering how to use the function.

    Parameters
    -----
    param1 : float
        Controls the adiabatic expansion of the universe. Use larger values for
        more excitement. Use smaller values if you're a wall liker. Use
        negative values if you're a maniac.

    param2 : int
        A flag which indicates what I ate for dinner last night, where 0 is
        Subway, 1 is chow mein, and 2 is curry. Values above 2 are never used.

    Other Parameters
    -----
    param3 : None, str, default = None
        If not None, then represents a love letter that will be sent to That
        Game Company for making Journey, the best game ever.



    Returns
    -----
    stuff : float
        The stuff that this function returns.
    """

    # param1 is mostly a joke
    if param1 < 0: print('What have you done...?')

    dinner = None
    if param2 == 0: dinner = 'Subway'
    elif param2 == 1: dinner = 'chow mein'
    elif param2 == 2: dinner = 'curry'
    else: raise NotImplementedError('Unrecognized value for param2: %d' % param2)




    if param3 is not None: # Only if something to send
        print("I hope this arrives at That Game Company HQ...\n" + param3)

    # Returning junk value because this function is a joke
    stuff = 1.
    return stuff
```



Thinking in parallel: mail delivery

- Deliver & pickup info
- You are the director @ HQ
- Serial (1 truck)
 - Simple
 - Takes too long
- Parallel (many trucks)
 - More complex
 - Much faster



- Example with single “child” process
- When we run the script, we create the main process
- Main creates a child process
- If main dies, the child should die too
- “daemon = True”
- This is still equivalent to serial

scripts/examples/single_process.py




```
import multiprocessing, time

def func(): time.sleep(1)


if __name__ == '__main__': # Only if we are the main process
    # Create a child process
    process = multiprocessing.Process(
        target = func,
        daemon = True, # Terminate child when main exits
    )
    process.start()
    start_time = time.time() # Starting timestamp
    process.join() # Wait until child finishes
    end_time = time.time() # Stopping timestamp
    print(end_time - start_time) # Expect: 1 (s)
```

How to direct to the mail drivers?


- Create a “queue” at HQ
 - Driver calls for new address: take from the queue
- multiprocessing.Queue
 - “First-in, first-out” (FIFO)
 - “Last-in, first-out” (LIFO)



3 houses, 3 drivers? EZ



uh-oh



GOAL:

```
$ python3 queues.py  
0  
1  
2  
3  
4  
5  
6  
7
```


1) Make Queue

```
import multiprocessing

if __name__ == '__main__':
    nprocs = 4
    queue = multiprocessing.Queue()

    # Put values in the queue
    for i in range(nprocs*2): queue.put(i)
    # Append end-of-queue signals
    for _ in range(nprocs): queue.put(None)

    # Create nprocs processes
    processes = [multiprocessing.Process(
        target = func,
        args = [queue],
        daemon = True,
    ) for _ in range(nprocs)]

    # Start processes
    for process in processes: process.start()
    # Wait for processes to finish
    for process in processes: process.join()
```

2) Fill Queue

When a process reads
None from the queue, it
stops running

```
import multiprocessing

if __name__ == '__main__':
    nprocs = 4
    queue = multiprocessing.Queue()

    # Put values in the queue
    for i in range(nprocs*2): queue.put(i)
    # Append end-of-queue signals
    for _ in range(nprocs): queue.put(None)

    # Create nprocs processes
    processes = [multiprocessing.Process(
        target = func,
        args = [queue],
        daemon = True,
    ) for _ in range(nprocs)]

    # Start processes
    for process in processes: process.start()
    # Wait for processes to finish
    for process in processes: process.join()
```

3) Create processes

```
import multiprocessing

if __name__ == '__main__':
    nprocs = 4
    queue = multiprocessing.Queue()

    # Put values in the queue
    for i in range(nprocs*2): queue.put(i)
    # Append end-of-queue signals
    for _ in range(nprocs): queue.put(None)

    # Create nprocs processes
    processes = [multiprocessing.Process(
        target = func,
        args = [queue],
        daemon = True,
    ) for _ in range(nprocs)]

    # Start processes
    for process in processes: process.start()
    # Wait for processes to finish
    for process in processes: process.join()
```

3) Create processes

```
def func(queue):  
    while True: # until no more tasks  
        i = queue.get()  
        if i is None: # end-of-queue  
            return  
        print(i)
```

```
import multiprocessing  
  
if __name__ == '__main__':  
    nprocs = 4  
    queue = multiprocessing.Queue()  
  
    # Put values in the queue  
    for i in range(nprocs*2): queue.put(i)  
    # Append end-of-queue signals  
    for _ in range(nprocs): queue.put(None)  
  
    # Create nprocs processes  
    processes = [multiprocessing.Process(  
        target = func,  
        args = [queue],  
        daemon = True,  
    ) for _ in range(nprocs)]  
  
    # Start processes  
    for process in processes: process.start()  
    # Wait for processes to finish  
    for process in processes: process.join()
```

4) Start processes

5) Wait

```
def func(queue):  
    while True: # until no more tasks  
        i = queue.get()  
        if i is None: # end-of-queue  
            return  
        print(i)
```

scripts/examples/queues.py

```
import multiprocessing  
  
if __name__ == '__main__':  
    nprocs = 4  
    queue = multiprocessing.Queue()  
  
    # Put values in the queue  
    for i in range(nprocs*2): queue.put(i)  
    # Append end-of-queue signals  
    for _ in range(nprocs): queue.put(None)  
  
    # Create nprocs processes  
    processes = [multiprocessing.Process(  
        target = func,  
        args = [queue],  
        daemon = True,  
    ) for _ in range(nprocs)]  
  
    # Start processes  
    for process in processes: process.start()  
    # Wait for processes to finish  
    for process in processes: process.join()
```

6) Output

```
def func(queue):  
    while True: # until no more tasks  
        i = queue.get()  
        if i is None: # end-of-queue  
            return  
        print(i)
```

```
$ python3 queues.py
```

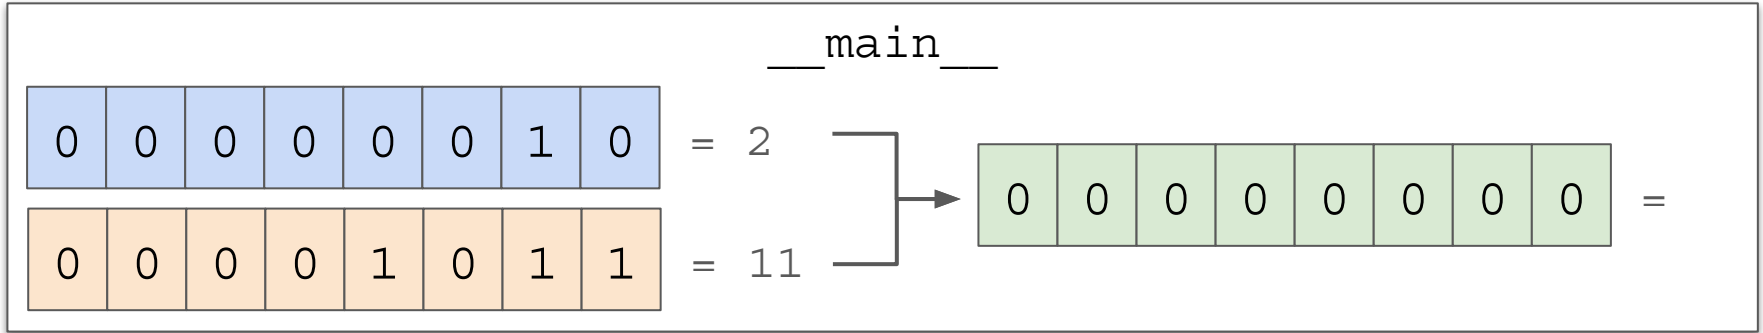
```
0  
1  
2  
3  
4  
5  
6  
7
```

scripts/examples/queues.py

```
import multiprocessing  
  
if __name__ == '__main__':  
    nprocs = 4  
    queue = multiprocessing.Queue()  
  
    # Put values in the queue  
    for i in range(nprocs*2): queue.put(i)  
    # Append end-of-queue signals  
    for _ in range(nprocs): queue.put(None)  
  
    # Create nprocs processes  
    processes = [multiprocessing.Process(  
        target = func,  
        args = [queue],  
        daemon = True,  
    ) for _ in range(nprocs)]  
  
    # Start processes  
    for process in processes: process.start()  
    # Wait for processes to finish  
    for process in processes: process.join()
```

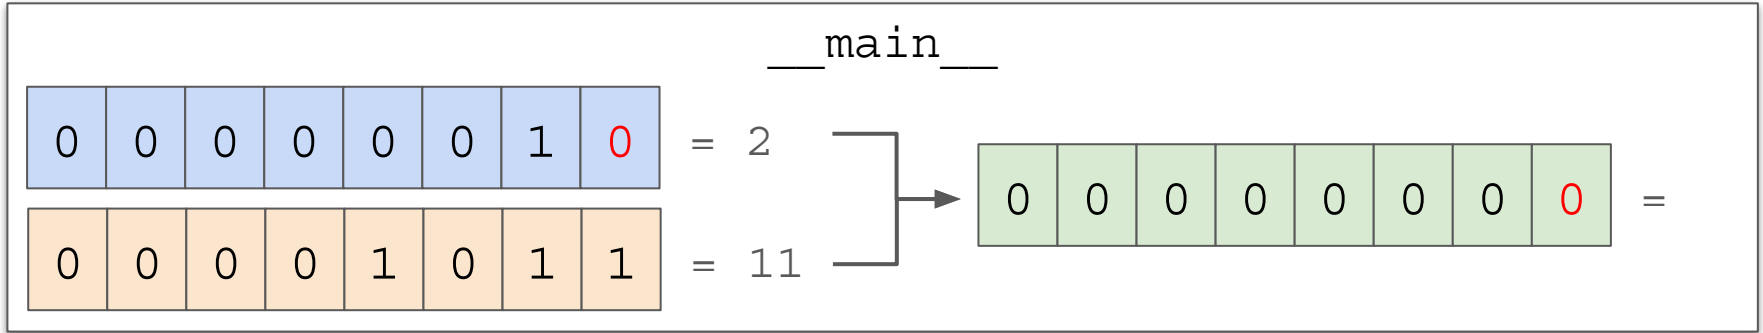
Complications

- Consider summation in **series**



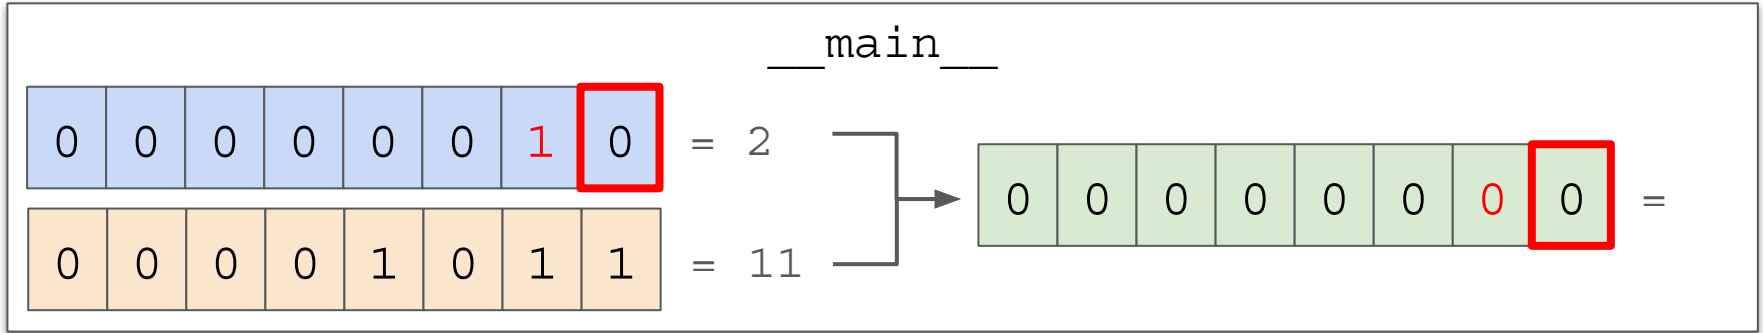
Complications

- Consider summation in **series**



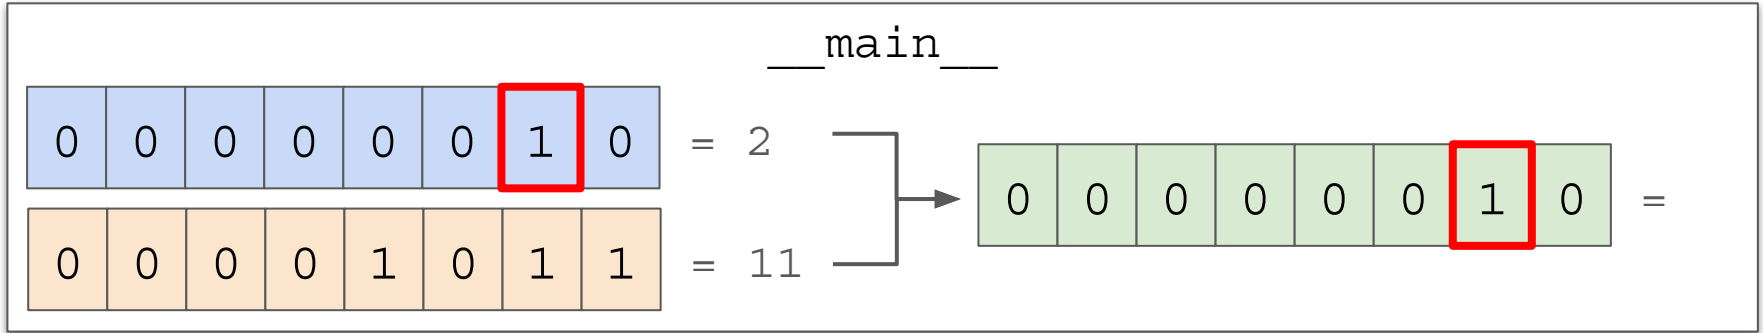
Complications

- Consider summation in **series**



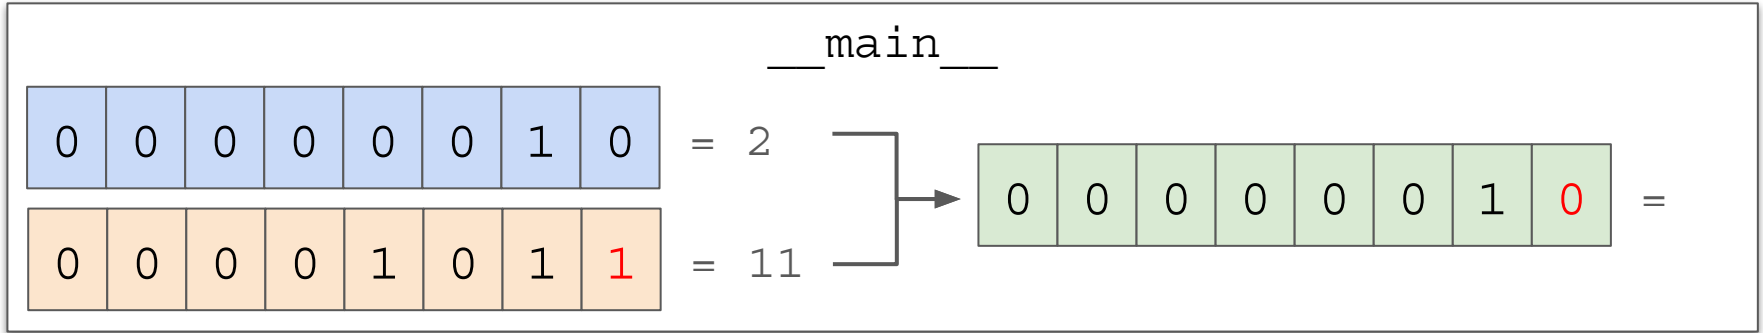
Complications

- Consider summation in **series**



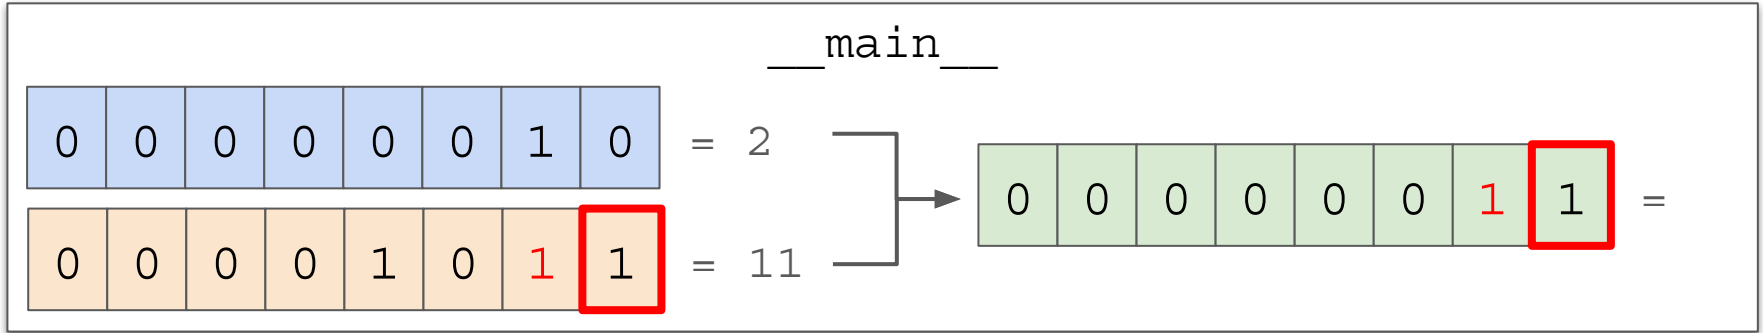
Complications

- Consider summation in **series**



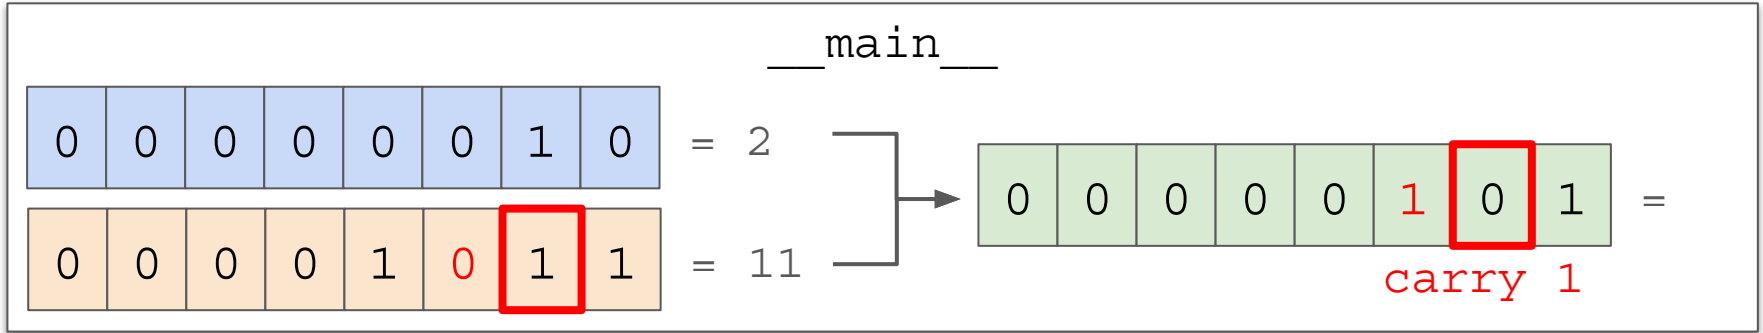
Complications

- Consider summation in **series**



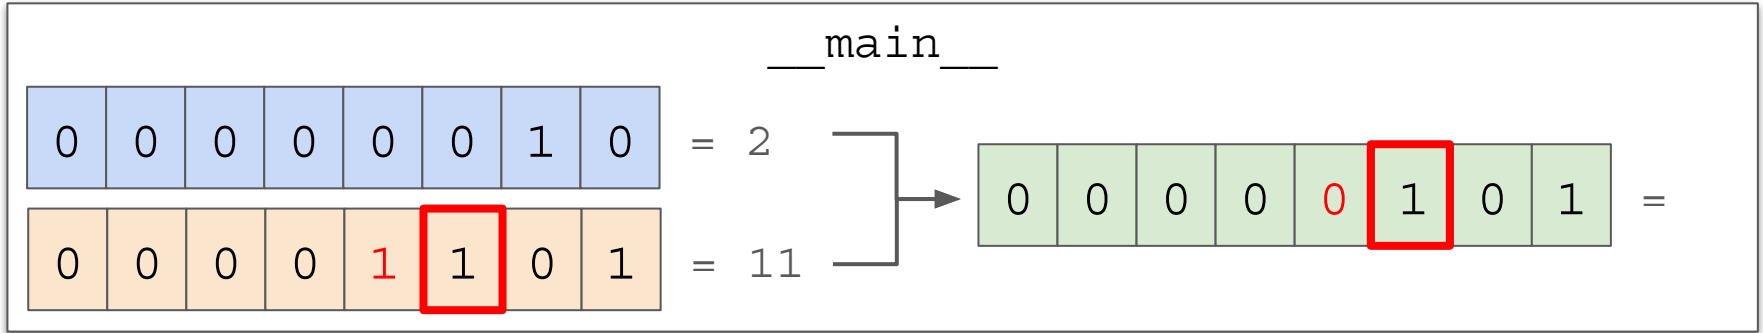
Complications

- Consider summation in **series**



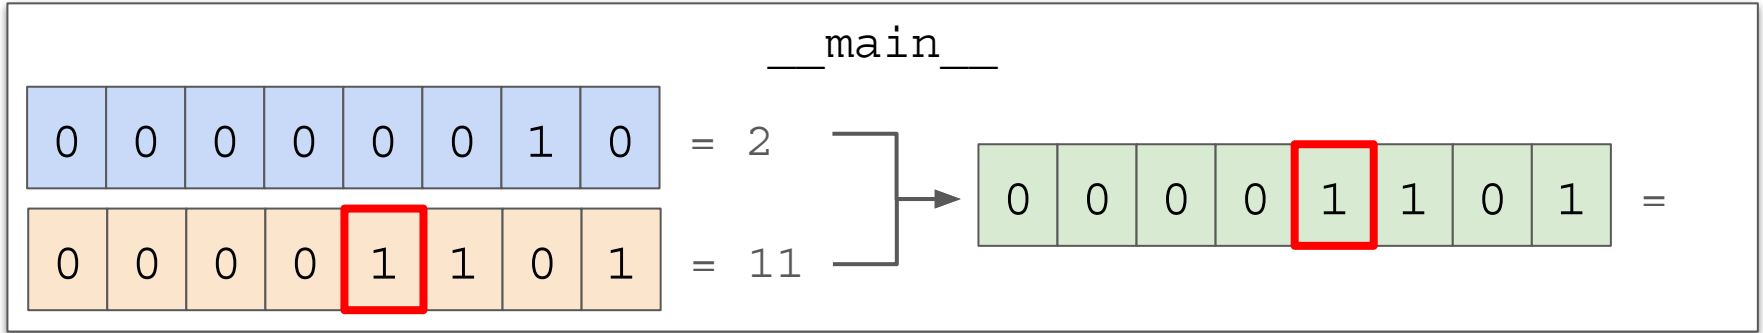
Complications

- Consider summation in **series**



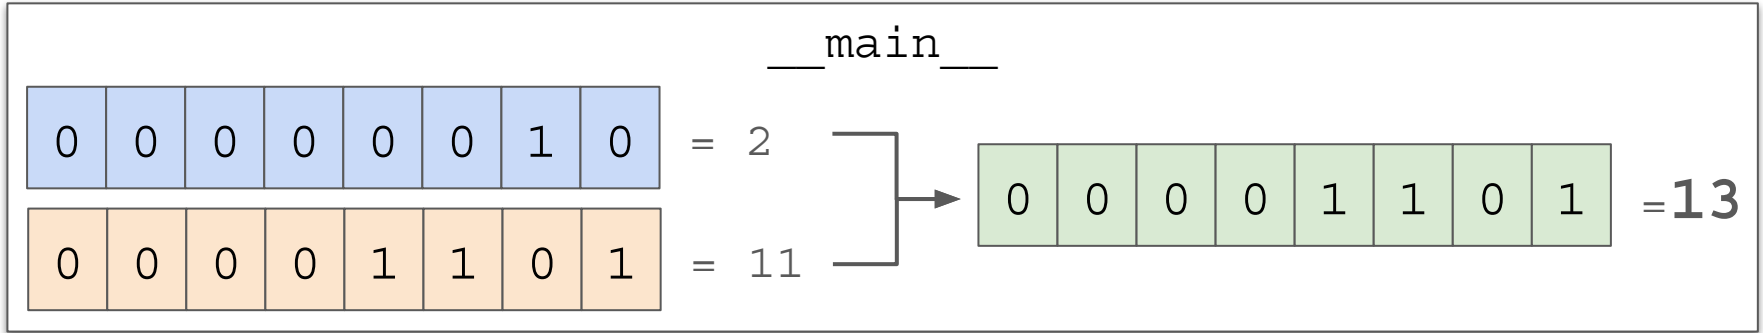
Complications

- Consider summation in **series**



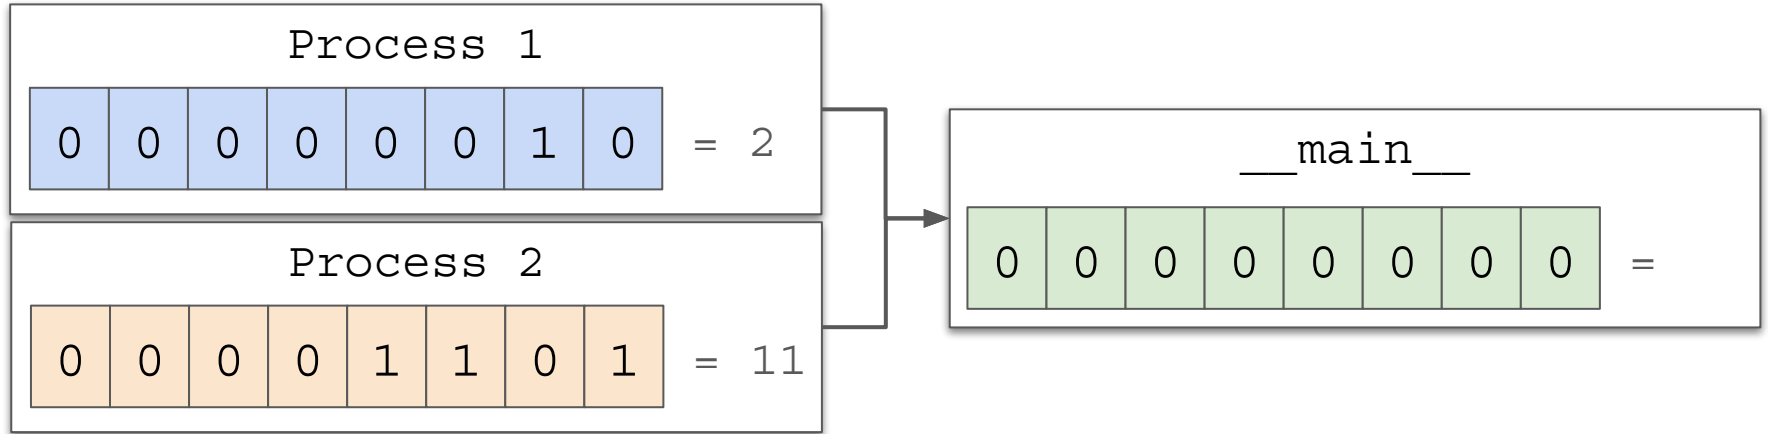
Complications

- Consider summation in **series**



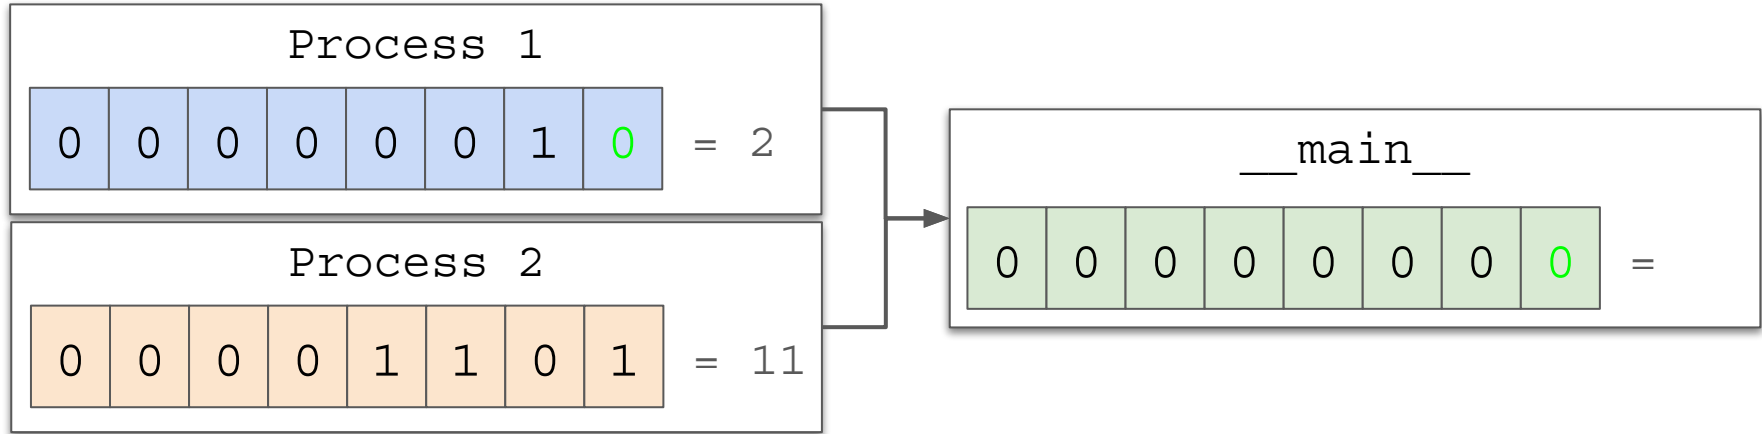
Complications

- Consider summation in **parallel**
 - But processes each write to same place in memory



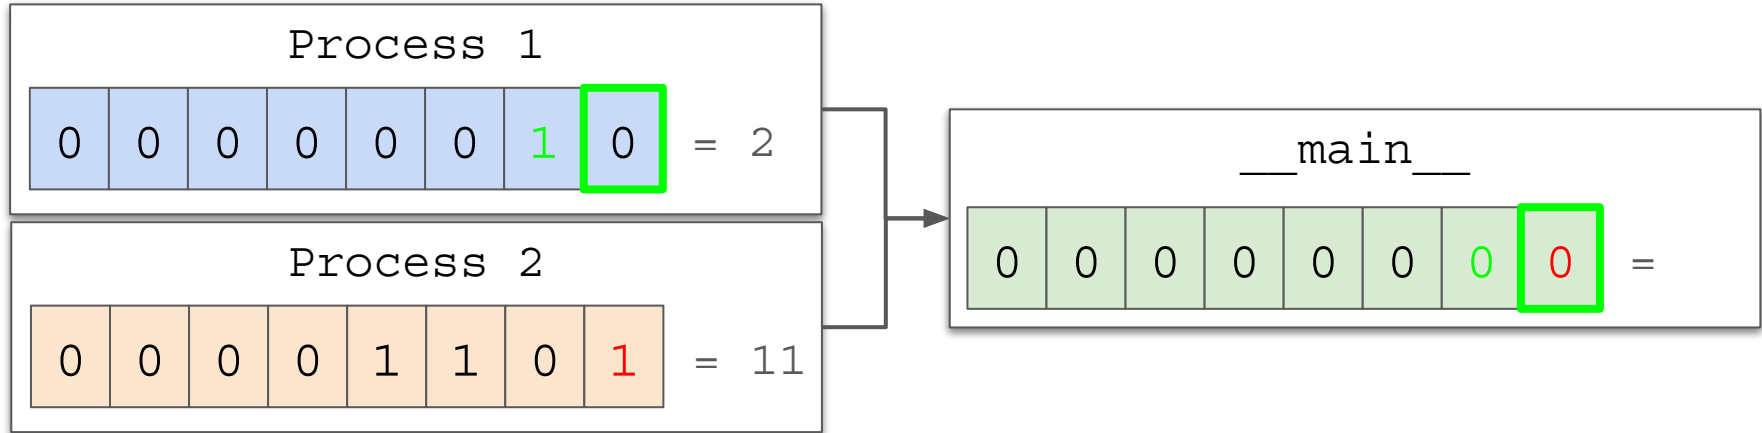
Complications

- Consider summation in **parallel**
 - But processes each write to same place in memory






Complications

- Consider summation in **parallel**
 - But processes each write to same place in memory






Complications

- Consider summation in **parallel**
 - But processes each write to same place in memory



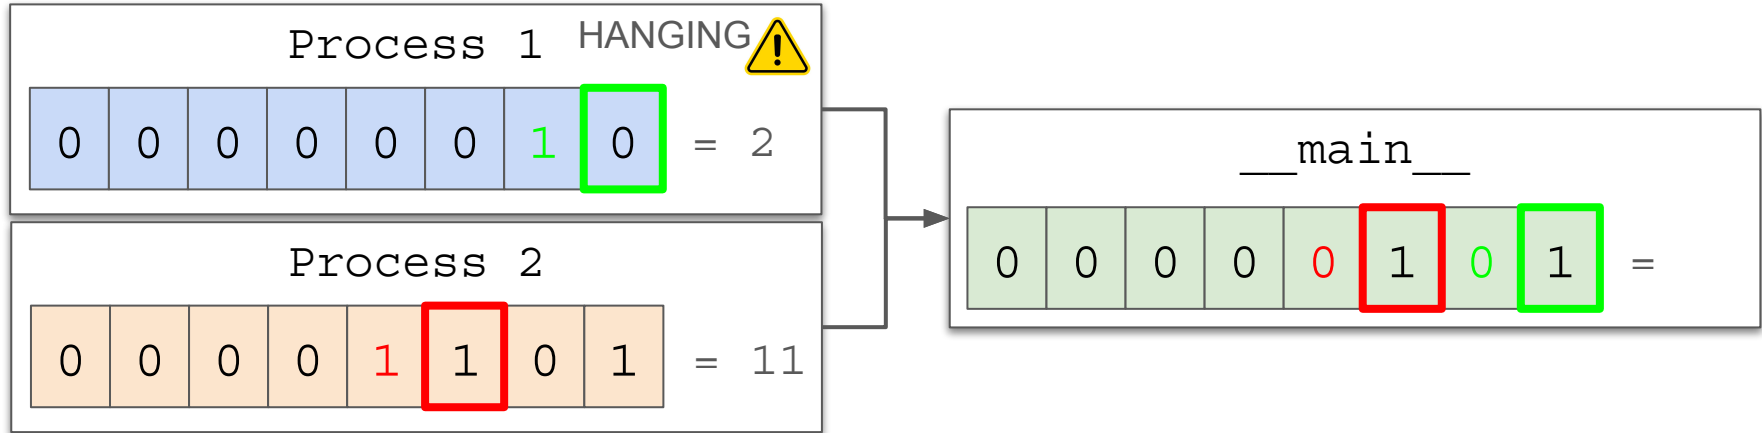
Complications

- Consider summation in **parallel**
 - But processes each write to same place in memory



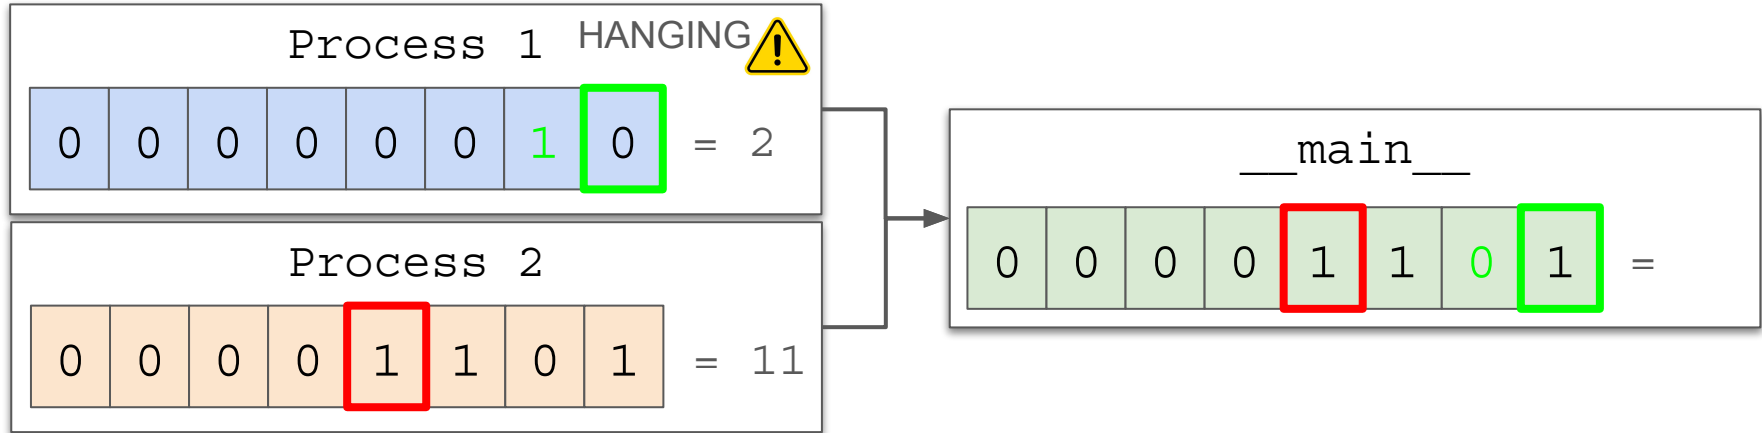
Complications

- Consider summation in **parallel**
 - But processes each write to same place in memory



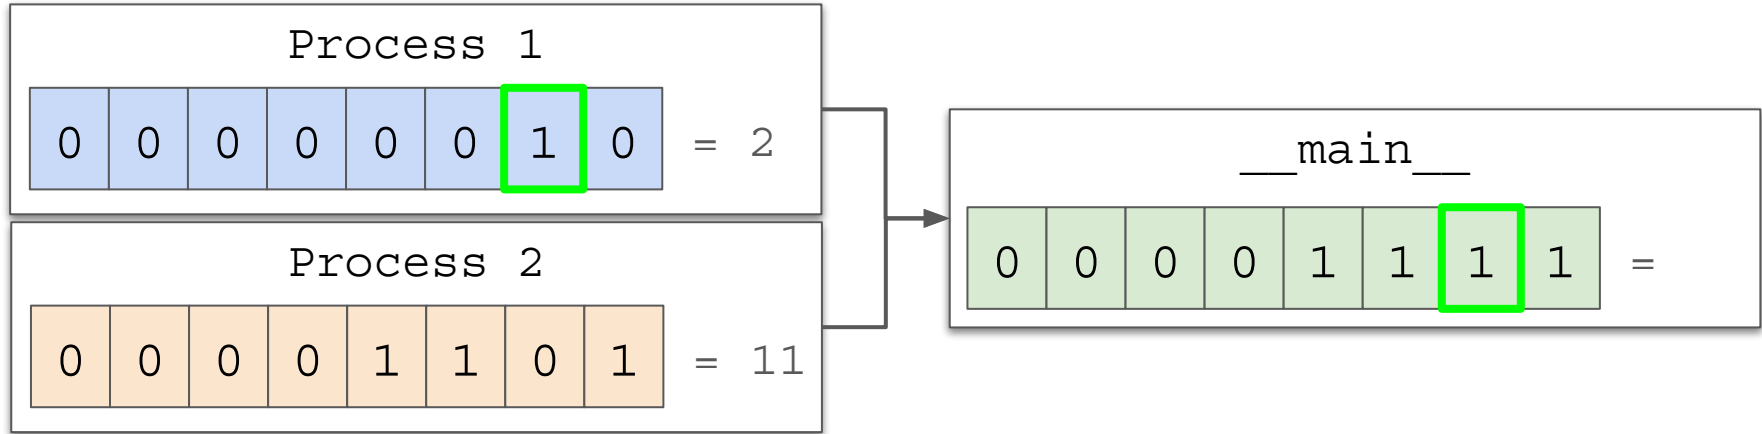
Complications

- Consider summation in **parallel**
 - But processes each write to same place in memory



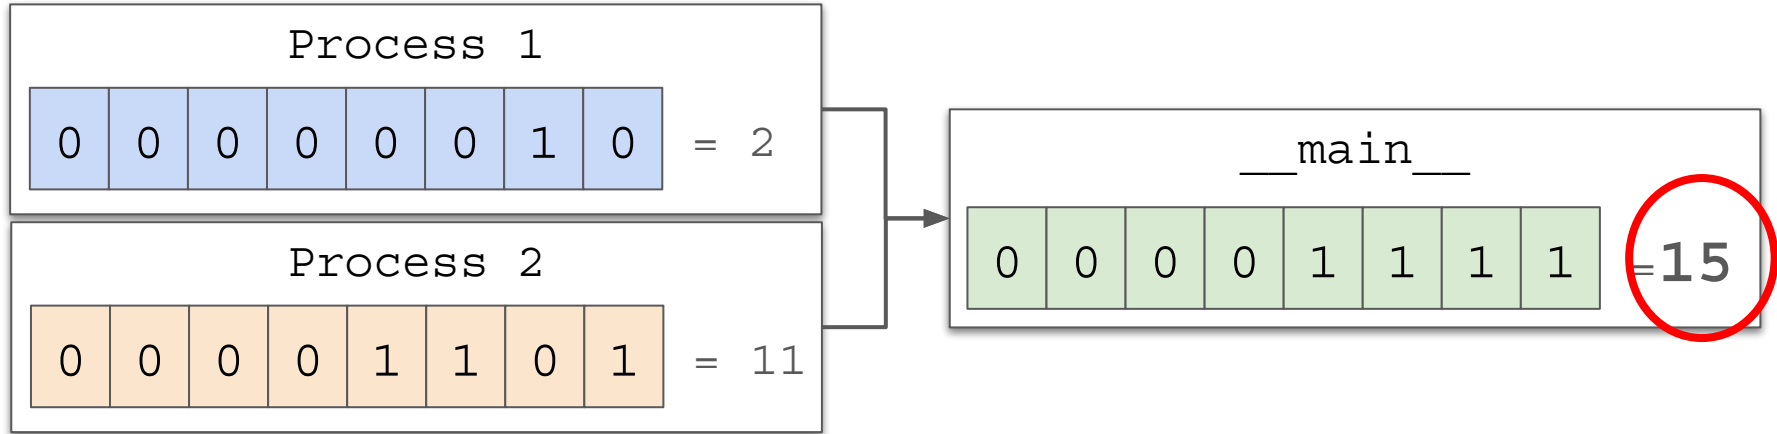
Complications

- Consider summation in **parallel**
 - But processes each write to same place in memory



Complications

- Consider summation in **parallel**
 - But processes each write to same place in memory



File Management

- Race conditions could fatally corrupt files
- BUT: Drives are serial
 - Except for niche cases like RAID
- 1 “head”: reads and writes

- How parallel can help:
 - Main process reads data and sends to processes
 - Processes put results in a queue
 - Main process retrieves results from queue
- Only helpful for long-running calculations

(snippet) `scripts/examples/fileio.py`


```
for process in processes: process.start()

# Fill the input queue
for line in read(filename): input_queue.put(line)
for _ in processes:
    input_queue.put(None) # end-of-queue signal

# Get the results from the output queue
ndone = 0
while ndone < len(processes):
    process_errors() # Check for errors
    if output_queue.empty(): continue
    if output_queue.get() is None: ndone += 1
```

With great power comes great code

- Parallel sum (“reduce”)
- Suppose we want: $\text{result} = 0 + 1 + 2 + \dots + n$
 - Say, $n = 2 \times 10^9$ 😈
 - In serial: `sum(range(int(2e9))) = 1999999999000000000`
- Try it yourself! `workspace/reduce.py`
 - Must use exactly 4 processes, and be able to prove it
 - Allowed modules: `multiprocessing`, `NumPy`, `time`, and `queue`
 - Cannot edit code marked with “DO NOT EDIT”
 - Local machines only
 - Everything else is allowed (go crazy)
 - **First code to execute in <1.55 seconds wins \$10**
 - Hints:
 - Remember: pre-fetching!
 - Allocating too much memory on Ubuntu makes your kernel crash



(github page)

My solution

- “Chunking”: using 10 chunks

```
def func(input_queue, output_queue):  
    while True:  
        try: output_queue.put(  
            np.einsum(  
                'i->',  
                np.arange(*input_queue.get()),  
            )  
        )  
    except (queue.Empty, TypeError):  
        output_queue.put(None)  
        return
```

```
$ python3 reduce.py  
That took 1.5385827461723238 seconds  
Correct result!
```

```
n = int(2e9)
```

```
nprocs = 4
```

```
input_queue = multiprocessing.Queue()
```

```
output_queue = multiprocessing.Queue()
```

```
chunklen = n // nchunks
```

```
for i in range(nchunks):
```

```
    input_queue.put([i*chunklen, (i+1)*chunklen])
```

```
for _ in range(nprocs): input_queue.put(None)
```

```
processes = [multiprocessing.Process(  
    target = func,
```

```
    args = (input_queue, output_queue),
```

```
    daemon = True,
```

```
) for _ in range(nprocs)]
```

```
for process in processes: process.start()
```

```
result = 0
```

```
ndone = 0
```

```
while ndone < len(processes):
```

```
    if output_queue.empty(): continue
```

```
    r = output_queue.get()
```

```
    if r is None:
```

```
        ndone += 1
```

```
        continue
```

```
    result += r
```